



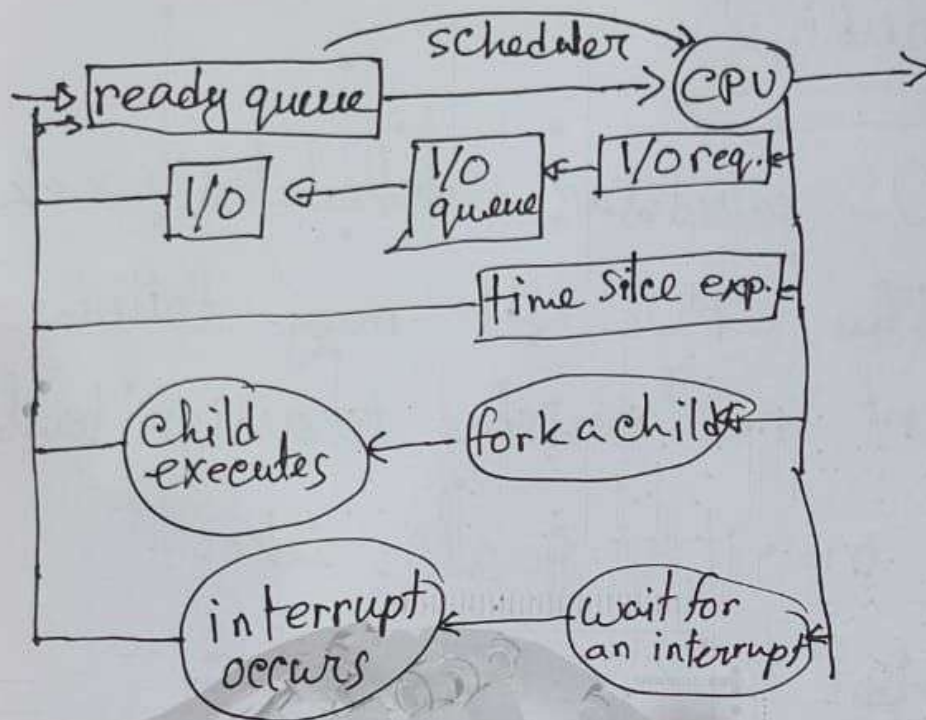
## \*Scheduling

⇒ An OS scheduler schedules work onto the CPUs. It may follow different mechanisms to assign tasks based on its goals, like:

- (i) Assign tasks immediately  
goal: - scheduling is simple (First come first serve (FCFS))
- (ii) Assigns simple tasks first  
o - maximize throughput (shortest job first (SJF))
- (iii) Assigns complex tasks first  
- maximize utilization of CPU, devices, memory.



## \* CPU Scheduling



→ chooses one of ready tasks to run on the CPU

→ scheduler runs when

⇒ CPU becomes idle

⇒ new tasks become ready

⇒ timeslice expired timeout.

→ Thread is dispatched on CPU.

↳ before: context switch, enter user mode, set PC & go!





icpc International College  
Programming Contest  
icpc foundation



In summary:

scheduling = choose tasks from queue.

→ which task should be selected?

→ depends on the scheduling policy / algorithm

⇒ How this is done?

→ dep. on the runqueue data-structure.

runqueue  $\Leftrightarrow$  scheduling algor.  
↳ tightly coupled

① "Run-to-completion" scheduling

This type of sch. assumes as soon as a task is assigned to a CPU it will run & finishes or completes.



Initial assumptions:

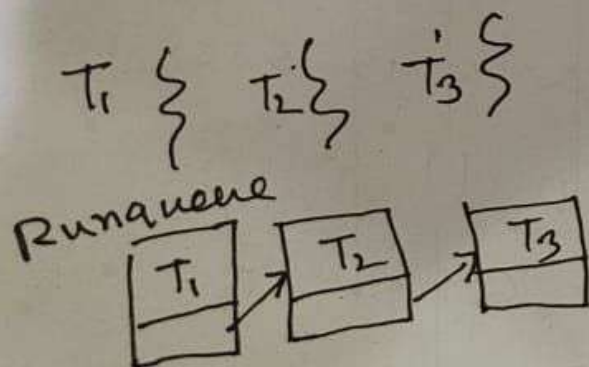
- group of tasks / jobs
- known execution times
- no preemption
- single CPU

metrics for comp algo:

- throughput,
- avg. job completion time
- avg. " waiting time
- CPU utilization

(1) First-Come First-Serve (FCFS)

- schedules tasks in order of arrival



runqueue ==  
queue (FIFO)



$$T_1 = 1s, T_2 = 10s, T_3 = 1s \dots (i)$$

Throughput:  $3/(1+10+1) = 0.25 \text{ tasks/s.}$

Avg. completion:  $(1+11+12)/3 = 8 \text{ sec}$

Avg. wait:  $(0+1+11)/3 = 4 \text{ sec.}$


poor if one  
task is long

## (ii) Shortest Job First (SJF)

→ Schedules tasks in order of  
their execution time

→ (i) →  $T_1(1s) \rightarrow T_3(1s) \rightarrow T_2(10s)$

runqueue = ordered queue.  
or tree DS.



runs the left first.

insertion  
is complex  
but  
retrieval  
is not



Throughput :  $(1+10+1)/3 = 0.25/s$


Avg. completion:  $(1+2+12)/3 = 5s$

Avg. wait:  $(0+1+2)/3 = 1$

\* Preemptive scheduling

SJF + Preemption

Task	Exec time	Arrival time
T1	1s	2
T2	10s	0
T3	1s	2



~~if we~~ In actual scenario, we do not know the exec. time





icpc International College  
Programming Contest  
icpc.foundation



BUBT



so, we need a heuristic based on  
history  $\Rightarrow$  job running time  
 $\rightarrow$  how long did a task run  
last time?  
 $\rightarrow$  how long did a task run last  
n time? (window avg).

### ③ Priority scheduling (a part of Preemptive sch)

- tasks have different priority levels.
- run highest priority runs next.

Task	Exec	arrival	Priority
$T_1$	1	2	$P_1$
$T_2$	10	0	$P_2$
$T_3$	1	2	$P_3$

$$P_1 < P_2 < P_3$$



→ running scenario:

0-2 →  $T_2$


2-3 →  $T_3$  (as this has the highest <sup>prio</sup>)

3-10 →  $T_2$  (second highest)

10-11 →  $T_1$  (as lowest)

runqueue for this is:

⇒ we can have multiple runqueues ds.



⇒ tree ordered based on Priority.

⇒ demerits:

starvation: as low priority task  
stuck in a run queue.

one possible solution: "priority

aging".  $p = f(\text{actual prio.}, \text{time spent in run queue})$ .

\*A math in the notes






icpc International College  
Programming Contest  
icpc.foundation



\*Priority Inversion: assumes SJF here.

Task	Arr. time	Priority
$T_1$	5	$P_1$
$T_2$	3	$P_2$
$T_3$	0	$P_3$

$P_1 > P_2 > P_3$



Order of execution:

$T_2, T_3, T_1$

⇒ Priorities "inverted".

Solution:

- temp boost priority of mutex owner.




## (iv) Round Robin Scheduling

- Pick up first task from queue.
- (FCFS)
- task may yield, to wait on I/O (unlike FCFS).

Task	Exec	Arr
T <sub>1</sub>	2	0
T <sub>2</sub>	2	0
T <sub>3</sub>	2	0

Runqueue



# more in notes for example

\* Timeslice: maximum amount of uninterrupted time given to a task → time quantum

→ task may run less than timeslice time

- has to wait on I/O, sync...  
⇒ will be placed on a queue






icpc International College  
Programming Contest  
icpc foundation



- higher priority task becomes runnable.
- using timeslices tasks are interleaved.  $\Rightarrow$  timesharing the CPU
- CPU bound task  $\rightarrow$  preempted after timeslice.



Throughput	throughput	Avg. wait	avg. com
Alg			
FCFS	0.25 t/s	4s	8s
SJF	0.25 t/s	1	5
RR( $t_s=1$ )	0.25 t/s	1s	5.33

- + short task finish sooner
- + more responsive
- + lengthy I/O ops initiated sooner
- Overhead
- @sol: keep  $t_s \gg$  context switch



12:12  
# How long should a timeslice be?

⇒ we have to keep in mind to balance benefits and overhead.

sc1: → for I/O bound

sc2 → for CPU "

# discuss these scenarios with example

ICPC  
Asia Dhaka  
Regional Contest  
October, 2022