


## DISTRIBUTED SYSTEMS (COMP9243)

### Lecture 9b: Middleware

#### Slide 1

- ① Introduction
- ② Distributed Object Middleware
  - Remote Objects & CORBA
  - Distributed Shared Objects & Globe
- ③ Publish/Subscribe Middleware

#### Slide 2




### KINDS OF MIDDLEWARE

Distributed Object based:

→ Objects invoke each other's methods

#### Slide 3

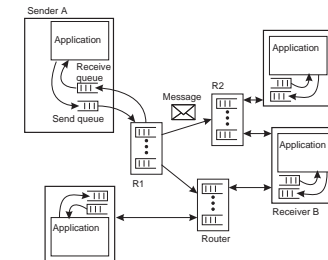


Message-oriented:

→ Messages are sent between processes

→ Message queues

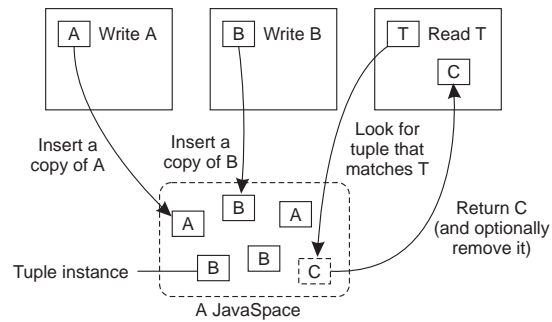
#### Slide 4



Slide 5

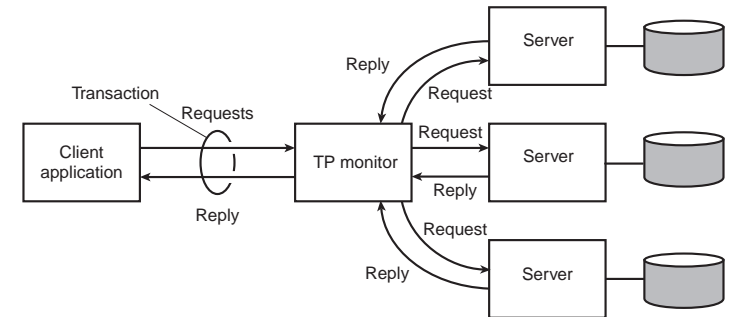
### Coordination-based:

→ Tuple space



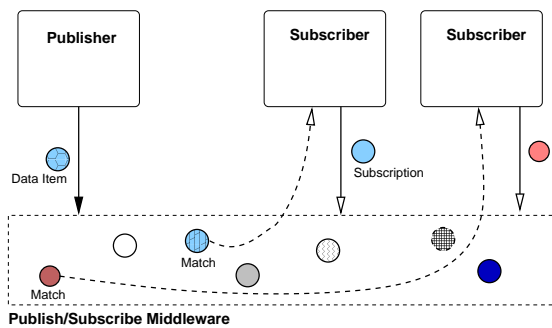
Slide 7

### Transaction Processing Monitors:




Slide 6

### → Publish/Subscribe




Slide 8

### Web Services:



---

## DISTRIBUTED OBJECTS



Slide 9

---

## CHALLENGES

- Transparency
    - Failure transparency
  - Reliability
    - Dealing with *partial failures*
  - Scalability
    - Number of clients of an object
    - Distance between client and object
  - Design
    - Must take distributed nature into account from beginning
  - Performance
  - Flexibility
- 

---

## OBJECT MODEL

- Classes and Objects
    - Class:** defines a type
    - Object:** instance of a class
  - Interfaces
  - Object references
  - Active vs Passive objects
  - Persistent vs Transient objects
  - Static vs Dynamic method invocation
- 

Slide 11

---

## INTERFACES

Define an object's methods


- Composition of interfaces
- Inheritance of interfaces
- Versions of interfaces

Slide 12

Interface Definition Language:

- Types of objects,
  - Public methods,
  - Exceptions that may occur
  - etc.
-

## REMOTE OBJECT ARCHITECTURAL MODEL



Slide 13

### Remote Objects:

- Single copy of object state (at single object server)
- All methods executed at single object server
- All clients access object through proxy
- Object's location is location of state

## CLIENT

### Client Process:


- Binds to distributed object
- Invokes methods on object

### Proxy:

- Proxy: RPC stub + destination details
- Binding causes a proxy to be created
- Responsible for marshaling
- Static vs dynamic proxies
- Usually generated

### Run-Time System:

- Provides services (translating references, etc.)
- Send and receive



Slide 14

## OBJECT SERVER

### Object:

- State & Methods
- Implements a particular interface

### Skeleton:


- Server stub
- Static vs dynamic skeletons

### Run-Time System:

- Dispatches to appropriate object
- Invocation policies

### Object Server:

- Hosts object implementations
- Transient vs Persistent objects
- Concurrent access
- Support legacy code




Slide 15

## OBJECT REFERENCE

### Local Reference:

- Language reference to proxy




Slide 16

## OBJECT REFERENCE

Remote Reference:


→ Server address + object ID

Slide 17




Slide 18

→ Reference to proxy code (e.g., URL) & init data



→ Object name (human friendly, object ID, etc.)

Slide 19



What are the drawbacks and/or benefits of each approach?

## BINDING AND NAME RESOLUTION

Name Resolution:

→ Name → remote reference

- Reference info contained in name (e.g., URL), or
- Naming service stores name to reference mappings

Slide 20

Binding:

→ Remote reference → local reference

- Create a proxy
- Connect proxy to object server

## REMOTE METHOD INVOCATION (RMI)

Standard invocation (synchronous):

- Client invokes method on proxy
- Proxy performs RPC to object server
- Skeleton at object server invokes method on object
- Object server may be required to create object first

Other invocations:

- Asynchronous invocations
- Persistent invocations
- Notifications and Callbacks

Slide 21

## INVOCATION SEMANTICS

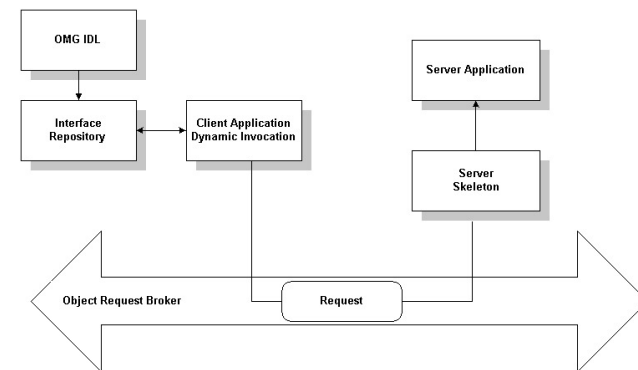
- Local method invocation: executed **exactly once**
- **Problem:** Cannot make such a guarantee for remote invocations!

Slide 22

Retransmit Request	Fault Tolerance Measure		Invocation Semantics
	Filter Duplicates	Re-execute, or Re-reply	
No	n/a	n/a	Maybe
Yes	No	Re-execute procedure	At-least-once
Yes	Yes	Retransmit reply	At-most-once

What's the difference between Maybe and At-most-once?

## STATIC VS DYNAMIC INVOCATION



Slide 23

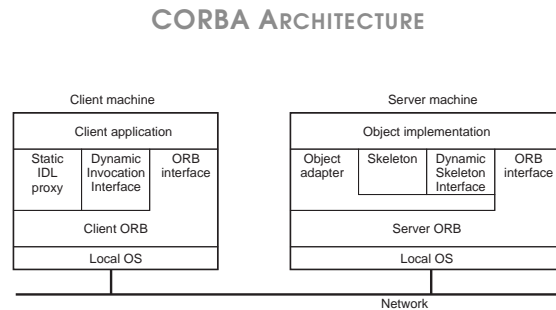
## CORBA

Features:

- Object Management Group (OMG) Standard (version 3.1)
- Range of language mappings
- Transparency: Location & some migration transparency
- Invocation semantics: at-most-once semantics by default; maybe semantics can be selected
- Services: include support for naming, security, events, persistent storage, transactions, etc.

Slide 24

Slide 25



Slide 27

### Example: A Simple File System:

```
module CorbaFS {
    interface File;          // forward declaration

    // a super-simple file system
    interface FileSystem {
        exception CantOpen {string reason;};
        enum OpenMode {Read, Write, ReadWrite};
        File open (in string fname, in OpenMode mode)
            raises (CantOpen);
    };

    // an open file
    interface File {
        string read (in long nchars);
        void write (in string data);
        void close ();
    };
};
```

Slide 26

### INTERFACES: OMG IDL

#### Components of an interface definition:

- Defines class attributes and methods  
(Classes are called **interfaces** and methods are called **operations**)
- Method arguments are annotated as **in**, **out**, and **inout**
- Errors/exceptions
- Interface inheritance

#### Characteristics of OMG IDL:

- Grammar is a subset of ANSI C++ plus additional constructs for object invocation
- No control structures, only declarations
- Primitive and complex data types, but **no pointers**
- No templates and no overloading
- Includes support for pre-processor

Slide 28

### OBJECT REFERENCE (OR)


#### Object Reference (OR):

- Refers to exactly one object, but an object can have multiple, distinct handles
- A language mapping provides an opaque representation
- ORs are implementation specific

Slide 29

## Interoperable Object Reference (IOR)

- Can be shared between different implementations




## CLIENT

- Binds to remote object
- Static proxy
- Dynamic invocation

### Client-Side Dynamic Invocation:


- Dynamic Invocation interface (DII)
- Request is dynamically created
- Interface repository:
  - Persistent storage of interface definitions
  - Dynamic type checking and checking of inheritance graph
  - Interface browser
  - May be queried by language bindings



Slide 30

## OBJECT

- Also known as *Servant*
- Can be proper object (e.g., C++, Java)
- Can be state and procedures (e.g., C)
- Can be legacy code (with a wrapper)



Slide 31

### Object Adapter:


- Dispatcher (OR → Servant)
- Invokes skeleton (static or dynamic)
- Possibly creates objects
- Portable Object Adapter

## OBJECT REQUEST BROKER (ORB)


- Provides run-time system
- Translate between remote and local references
- Send and receive messages
- Maintains interface repository
- Enables dynamic invocation (client and server side)
- Locates services

Slide 32





Slide 33



Slide 34

### CORBA Leaves Implementations a Lot of Freedom:

#### Advantages

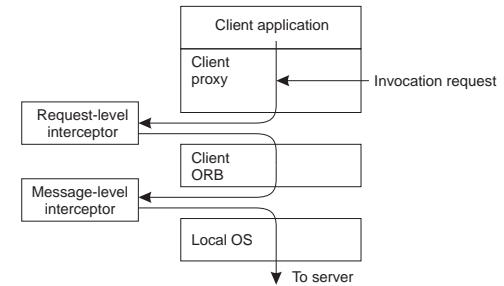
- ORB implementations can more easily be optimised for size, speed, or functionality
- Wide applicability; allows extreme cases like use in real-time systems (TAOS)
- Future improvements in network & compiler technology can be exploited more easily
- Vendors can specialise

#### Disadvantages

- Complexity of the specification; more difficult to understand
- Incompatibility of different implementations

This is in contrast to COM/DCOM.

### INTERCEPTORS



Slide 35


### BINDING

#### Direct Binding:

- Create proxy
- ORB connects to server (using info from IOR)
- Invocation requests are sent over connection

#### Indirect Binding:

Slide 36



**Slide 37**

- ## Components of the GIOP specification:

- ## Goals of GIOP:

- ## COMMON DATA REPRESENTATION (CDR) ✦

- For example:

## Slide 38

- ## Slide 40

- 20

## TRANSPORT LAYER (GIOP REQUIREMENTS)

- Transport must be connection-oriented
- Transport must be reliable
- Unbounded stream of bytes
- Transport must notify in case of connection loss
- TCP's connection initiation model must be implementable

Slide 41

→ TCP/IP fits very well

From GIOP to Internet Inter-Orb Protocol (IIOP):

A small step:

- IOR definition for TCP/IP
- Connection handling

## CORBA SERVICES

Some of the standardised services are the following:

- Naming Service
- Event Service
- Transaction Service
- Security Service
- Fault Tolerance

Slide 42

## CORBA BIBLIOGRAPHY

(1) *IIOP Complete*, W. Ruh, T. Herron, and P. Klinker, Addison Wesley, 1999.

(2) *The Common Object Request Broker: Architecture and Specification (2.3.1)*, Object Management Group, 1999.


(3) *C Language Mapping Specification*, Object Management Group, 1999.

(4) *CORBAservices: Common Object Services Specification*, Object Management Group, 1998.

Play with CORBA. Many implementations available, including ORBit: <http://www.gnome.org/projects/ORBit2/>

Slide 43

## DISTRIBUTED SHARED OBJECT (DSO) MODEL




Slide 44

Distributed Shared Objects:

- Object state can be replicated (at multiple object servers)
- Object state can be partitioned
- Methods executed at some or all replicas
- Object location no longer clearly defined

## CLIENT


- Client has local representative (LR) in its address space
- Stateless LR
  - Equivalent to proxy
  - Methods executed remotely
- Statefull LR
  - Full state
  - Partial state
  - Methods (possibly) executed locally



Slide 45

## OBJECT SERVER

- Server dedicated to hosting LRs
- Provides resources (network, disk, etc.)
- Static vs Dynamic LR support
- Transient vs Persistent LRs
- Security mechanisms




Slide 47

### Location of LRs:

- LR only hosted by clients
- Statefull LR only hosted by object servers
- Statefull LRs on both clients and object servers

## OBJECT



Slide 46

## GLOBE (GLOBAL OBJECT BASED ENVIRONMENT)

### Scalable wide-area distributed system:

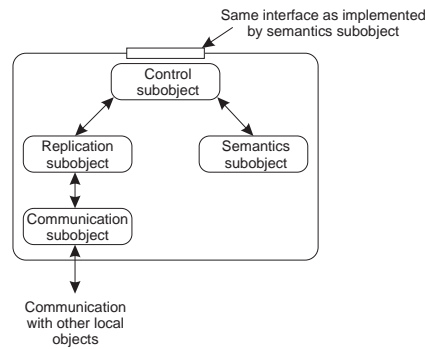
- Wide-area scalability requires replication
- Wide-area scalability requires flexibility

Slide 48

### Features:

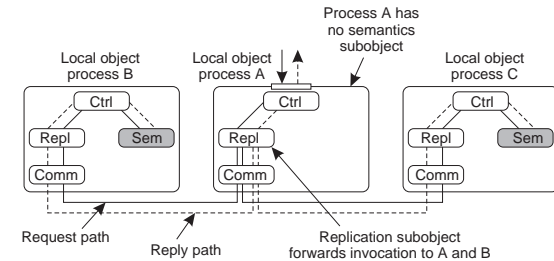
- Per-object replication and consistency
- Per-object communication
- Mechanism not policy
- Transparency (replication, migration)
- Dynamic replication

## LOCAL REPRESENTATIVE



Slide 49

## REPLICATION EXAMPLE: ACTIVE REPLICATION




Slide 51

### Semantics Subobject:

- Equivalent to CORBA Servant
- Implements object interface (Globe IDL)
- State stored in semantics subobject

### Control Subobject:

- Implements standard *control interface*
- Mediates between replication and semantics subobjects
- Does marshaling/unmarshaling

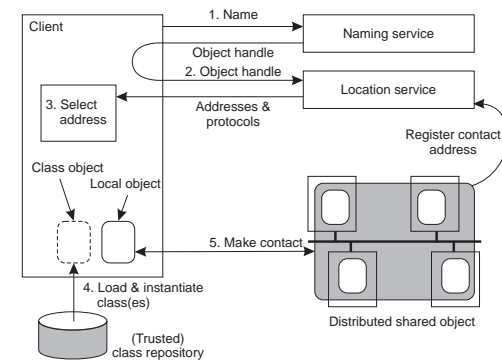


Slide 50

### Replication Subobject:

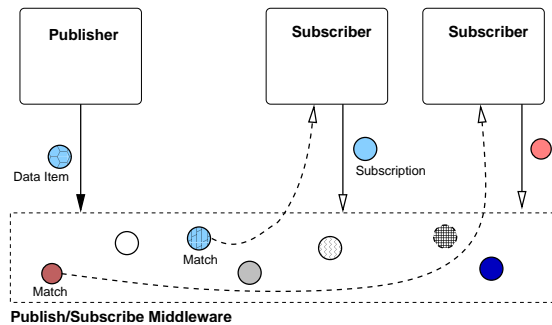
- Implements standard *replication interface*: `start()`, `send()`, `invoked()`
- Responsible for replication and consistency
- Implementation independent of semantics subobject

## BINDING



Slide 52

## PUBLISH/SUBSCRIBE (EVENT-BASED) MIDDLEWARE



Slide 53

## CHALLENGES

### Transparency:

- loose coupling → good transparency

### Scalability:

- Potentially good due to loose coupling
- ✗ In practice hard to achieve
- Number of subscriptions
- Number of messages

Slide 54

### Flexibility:

- Loose coupling gives good flexibility
- Language & platform independence
- Policy separate from mechanism

### Programmability:

- Inherent distributed design
- Doesn't use non-distributed concepts

## EXAMPLES

### Real-time Control Systems:

- External events (e.g. sensors)
- Event monitors

### Stock Market Monitoring:

- Stock updates
- Traders subscribed to updates

### Network Monitoring:

- Status logged by routers, servers
- Monitors screen for failures, intrusion attempts


### Enterprise Application Integration:

- Independent applications
- Produce output as events
- Consume events as input
- Decoupled

Slide 55


## MESSAGE FILTERING

### Topic-based




Slide 56

### Content-based




## ARCHITECTURE


Centralised:



Peer-to-Peer:



Multicast-based:



Slide 57

## COMMUNICATION

- Point-to-point
- Multicast
  - hard part is building appropriate multicast tree
- Content-based routing
  - point-to-point based router network
  - make forwarding decisions based on message content
  - store subscription info at router nodes

Slide 58

## REPLICATION

Replicated Brokers:

- Copy subscription info on all nodes
- Keep nodes consistent
- What level of consistency is needed?
- Avoid sending redundant subscription update messages

Slide 59

Partitioned Brokers:

- Different subscription info on different nodes
- Events have to travel through all nodes
- Route events to nodes that contain their subscriptions

## FAULT TOLERANCE

Reliable Communication:

- Reliable multicast

Process Resilience (Broker):

- Process groups
- Active replication by subscribing to group messages

Slide 60

Routing:

- Stabilise routing if a broker crashes
- Lease entries in routing tables

---

## EXAMPLE SYSTEMS

TIB/Rendezvous:

- Topic-based
- Multicast-based

Java Message Service (JMS):

- API for MOM
- Topic-based
- centralised or peer-to-peer implementations possible

Scribe:

- Topic-based
  - Peer-to-peer architecture, based on Pastry (DHT)
  - Topics have unique IDs and map onto nodes
  - Multicast for sending events
    - Tree is built up as nodes subscribe
- 
- 

Slide 61

## READING LIST

**Globe: A Wide-Area Distributed System** An overview of  
Globe

**CORBA: Integrating Diverse Applications Within Distributed  
Heterogeneous Environments** An overview of CORBA

**New Features for CORBA 3.0** More CORBA

---

Slide 62